

1. ◆ Handling Large Datasets in Power BI Online Service

Power BI Service supports different dataset storage & connection modes:

1. Import Mode

- Data is compressed and stored in memory (VertiPaq engine).
- Very fast for queries, but limited by dataset size quotas.
- Pro license → max **1 GB per dataset**.

2. DirectQuery Mode

- Data is not stored in Power BI; queries are sent to the source (SQL DB, etc.).
- No size limits, but performance depends on the source system.

3. Hybrid (Composite Models)

- Mix of Import + DirectQuery for balancing speed & size.
-

◆ Limitations with Pro License

- Dataset size limit: **1 GB per dataset**.
 - Refresh: up to **8 times/day**.
 - Shared cloud capacity → performance can be inconsistent.
-

◆ Role of Premium Capacity

Power BI Premium (P SKU for orgs, or PPU for individuals) provides **dedicated resources** in the Power BI Service.

Key Benefits:

1. Larger Datasets

- Premium supports **up to 400 GB per dataset** (vs 1 GB in Pro).
- Compression means this can represent **trillions of rows** in reality.

2. More Frequent Refresh

- Up to **48 refreshes/day** (every 30 minutes).
- Pro only allows 8/day.

3. Dedicated Performance

- Your reports run on **dedicated cloud capacity**, not shared with others.

- Consistent speed, better for enterprise users.

4. Advanced Features

- AI capabilities (AutoML, cognitive services).
- Paginated reports.
- Deployment pipelines (Dev → Test → Prod).
- Multi-Geo support (store data in different regions).

5. Unlimited Distribution

- With Premium, free users in your org can view shared reports (no Pro license needed).
-

◆ Example Scenario

- A global retailer wants to analyze **5 years of sales data (~200 GB)**.
 - With Pro: Impossible (limit 1 GB).
 - With Premium: Possible (400 GB limit + compression).
 - They can refresh every 30 minutes, keeping near real-time dashboards.
-

✓ In short:

Power BI Online Service handles large datasets using Import, DirectQuery, and hybrid models. But for **true enterprise scale** (bigger than 1 GB, high refresh frequency, consistent performance), you need **Premium Capacity**, which unlocks **larger dataset sizes, dedicated resources, and enterprise-only features**.

2. ◆ Power BI Data Connection Modes in the Service

1. Import Mode

- **How it works:** Data is copied and compressed into Power BI's in-memory engine (VertiPaq).
- **Performance:** Very fast (data is already in memory).
- **Refresh:** Data becomes stale unless refreshed (up to 8/day with Pro, 48/day with Premium).
- **Storage:** Counts toward dataset size quota (1 GB Pro, 400 GB Premium).

✓ Best for:

- Small to medium datasets

- When you need **fast performance**
- When source system is slow/unavailable for real-time queries

✗ Limitations:

- Refresh needed
 - Dataset size limits apply
-

2. DirectQuery Mode

- **How it works:** Data stays in the source; Power BI sends queries **live** at runtime.
- **Performance:** Depends on source system (can be slow if the DB isn't optimized).
- **Refresh:** Not needed (data always current), but metadata can still refresh.
- **Storage:** Very little space in Power BI since data isn't imported.

✓ Best for:

- Large datasets (can't fit into Import mode)
- Real-time reporting needs

✗ Limitations:

- Slower than Import (query round trips)
 - Limited DAX functionality (not all functions supported)
 - Puts load on the source system
-

3. Live Connection

- **How it works:** Similar to DirectQuery, but connects specifically to **Analysis Services models (SSAS, AAS, or existing Power BI datasets)**.
- **Data storage:** No data is stored in Power BI; it just queries the model directly.
- **Modeling:** Can't create new tables/relationships in Power BI Desktop — must use the model as-is.

✓ Best for:

- Centralized, curated **semantic models** (e.g., corporate data warehouse in SSAS).
- Consistent “single source of truth” across reports.

Limitations:

- No data mashup (only 1 live source per report).
 - Modeling restricted to the source.
 - Performance depends on SSAS/Power BI dataset performance.
-

◆ Quick Comparison Table

Feature	Import	DirectQuery	Live Connection
Data Storage	In Power BI (VertiPaq)	In source	In source (SSAS/PBI dataset)
Performance	 Very fast	▲ Depends on source	▲ Depends on source model
Refresh Needed	Yes (8/day Pro, 48 Premium)	No (always live)	No (always live)
Dataset Size Limit	1 GB (Pro), 400 GB (Premium)	None	None
DAX Support	Full	Limited	Limited
Modeling Flexibility	High	Medium	Low (no new tables/relations)
Best For	Small/medium data, fast reporting	Huge datasets, real-time	Centralized enterprise models

In short:

- **Import** = Fastest, but limited by dataset size & refresh.
- **DirectQuery** = No size limits, real-time, but slower & limited.
- **Live Connection** = Centralized corporate models, but very restrictive for modeling.

3. ◆ What are Deployment Pipelines?

Deployment pipelines help you **manage, test, and release Power BI content** in a structured way.

- Think of them as **Dev → Test → Prod** environments for Power BI.

- They ensure changes are validated before reaching business users.
-

- ◆ **Stages in a Deployment Pipeline**

1. **Development (Dev)** 

- Workspace for developers to build and experiment with reports, datasets, and dashboards.
- Frequent changes are expected.
- Only the dev team usually has access.

2. **Test (or QA)** 

- Workspace for testing and validation.
- Business analysts, testers, or selected users verify:
 - Data accuracy
 - Visual correctness
 - RLS rules
 - Performance
- Ensures everything works as expected before rollout.

3. **Production (Prod)** 

- Workspace for end users (read-only access).
 - Content here is **stable, validated, and approved**.
 - Distributed via **Apps** for broad usage across the organization.
-

- ◆ **Key Features of Deployment Pipelines**

- **Content Promotion:** Move content from Dev → Test → Prod with one click.
- **Change Detection:** Shows differences between stages (e.g., new reports, modified datasets).
- **Selective Deployment:** You can choose what to promote (not everything at once).
- **Data Source Rules:** You can configure connections differently per stage.
 - Example:
 - Dev → Test DB

- Test → QA DB
 - Prod → Production DB
-

◆ Benefits for Enterprises

- ✓ Consistency — same pipeline structure across projects
 - ✓ Governance — control who can publish at each stage
 - ✓ Faster releases — smooth transitions instead of manual uploads
 - ✓ Reduced risk — test before production rollout
-

✓ In short:

Deployment pipelines in Power BI Online give you **Dev → Test → Prod** stages for managing datasets, reports, and dashboards. They allow controlled, secure, and consistent promotion of content in enterprise environments.

4. ◆ Power BI + Microsoft Teams Integration

Power BI and Teams work together so reports are available right where people chat and collaborate.

1. Power BI App inside Teams

- You can install the **Power BI app** directly in Teams.
- This lets you:
 - Browse and search reports without leaving Teams
 - Pin important reports to the Teams left navigation
 - Access shared workspaces

2. Embed Reports in Teams Channels/Chats

- In a Teams channel → click + **(Add a tab)** → choose **Power BI** → select a report.
- Members can view and interact with the report without leaving Teams.
- Example: A Sales team can view a live Sales dashboard right in their Teams channel.

3. Chat about Data

- You can copy a link to a Power BI visual/report and paste it in Teams chat.
- The report opens interactively for people with permissions.

4. Notifications in Teams

- Using **Power Automate**, you can send alerts from Power BI to Teams when KPIs are reached or thresholds are broken.
-

◆ Power BI + SharePoint Online Integration

1. Embed Reports in SharePoint Pages

- Use the **Power BI web part** in SharePoint Online.
- Insert a live interactive Power BI report into a SharePoint modern page.
- No need to export static charts — users see real-time reports.

2. Permissions Handling

- Report permissions in Power BI **must match** SharePoint access.
- If someone doesn't have access in Power BI, embedding in SharePoint won't bypass security.

3. Use Cases

- Internal portals (e.g., HR dashboards, project updates, finance KPIs)
 - SharePoint becomes a central hub for both documents and live analytics.
-

✓ Summary

- **Teams** → best for **daily collaboration** (chat, channels, tabs with live reports).
- **SharePoint** → best for **company portals** (embed live reports in intranet pages).

Both integrations ensure that **data meets people where they already work**, instead of forcing them to open Power BI separately.

5. ◆ Power BI + Microsoft Teams Integration

Power BI and Teams work together so reports are available right where people chat and collaborate.

1. Power BI App inside Teams

- You can install the **Power BI app** directly in Teams.
- This lets you:
 - Browse and search reports without leaving Teams

- Pin important reports to the Teams left navigation
- Access shared workspaces

2. Embed Reports in Teams Channels/Chats

- In a Teams channel → click + **(Add a tab)** → choose **Power BI** → select a report.
- Members can view and interact with the report without leaving Teams.
- Example: A Sales team can view a live Sales dashboard right in their Teams channel.

3. Chat about Data

- You can copy a link to a Power BI visual/report and paste it in Teams chat.
- The report opens interactively for people with permissions.

4. Notifications in Teams

- Using **Power Automate**, you can send alerts from Power BI to Teams when KPIs are reached or thresholds are broken.

◆ Power BI + SharePoint Online Integration

1. Embed Reports in SharePoint Pages

- Use the **Power BI web part** in SharePoint Online.
- Insert a live interactive Power BI report into a SharePoint modern page.
- No need to export static charts — users see real-time reports.

2. Permissions Handling

- Report permissions in Power BI **must match** SharePoint access.
- If someone doesn't have access in Power BI, embedding in SharePoint won't bypass security.

3. Use Cases

- Internal portals (e.g., HR dashboards, project updates, finance KPIs)
- SharePoint becomes a central hub for both documents and live analytics.

✓ Summary

- **Teams** → best for **daily collaboration** (chat, channels, tabs with live reports).

- **SharePoint** → best for **company portals** (embed live reports in intranet pages).

Both integrations ensure that **data meets people where they already work**, instead of forcing them to open Power BI separately.

5. ◆ What is the XMLA Endpoint?

- **XMLA (XML for Analysis)** is a protocol that allows external tools to connect to **tabular models** (semantic models/datasets) in Power BI.
- Available only with **Premium capacity** (and Power BI Premium Per User).
- In practice, it makes a Power BI dataset behave like an **Analysis Services (SSAS) Tabular model** in the cloud.

The XMLA endpoint exposes your **Power BI datasets** for **read and write** operations depending on configuration:

- **Read-only mode:** Connect and query the dataset (e.g., with SSMS, Excel, or 3rd-party BI tools).
 - **Read/Write mode:** Advanced tasks like automation, metadata editing, partitions management, calculation groups, etc.
-

◆ Benefits for Developers & Enterprise BI Teams

1. Integration with External Tools

- Connect from **SQL Server Management Studio (SSMS), DAX Studio, Excel, or 3rd-party BI tools**.
- Run **DMVs** (Dynamic Management Views) to analyze dataset performance, model size, and queries.
- Treat Power BI like an **enterprise-grade OLAP engine**.

2. Advanced Modeling & Automation

- Manage partitions (useful for **incremental refresh** beyond UI limitations).
- Apply **tabular editor scripts** for calculation groups, measures, and metadata updates.
- CI/CD pipelines: Automate deployment of dataset changes via **DevOps**.

3. Enterprise-Scale Semantic Layer

- Multiple BI tools (Power BI, Excel, Tableau, etc.) can query the same **certified dataset** via XMLA.
- Promotes “**single source of truth**” for KPIs across the organization.

4. Performance Tuning

- Developers can capture and analyze query performance using **SQL Profiler** through XMLA.
- Helps optimize DAX queries, storage engine vs. formula engine bottlenecks.

5. Interoperability with Analysis Services

- Makes migration between **SSAS Tabular** and **Power BI Premium** easier.
 - Enterprises can treat Power BI Premium as a **cloud SSAS replacement**.
-

◆ Example Use Cases

- Finance team runs **Excel pivot tables** directly against a large certified Power BI dataset.
 - BI developers use **Tabular Editor** via XMLA to push advanced calculations.
 - DevOps pipelines push metadata/model changes automatically from Git → Premium dataset.
 - IT team monitors query performance using SSMS.
-

In short:

The **XMLA endpoint** turns Power BI Premium into a **full-fledged enterprise semantic model platform**, enabling advanced management, integration, automation, and scalability.

6. ◆ 1. Usage Metrics

- **What it is:** Built-in Power BI feature that shows **report and dashboard usage** (how many views, by whom, when).
- **Scope:** Per report or dashboard.
- **Where to find it:**
 - In Power BI Service → Open a report or dashboard → **More options (...)** → **Usage metrics**.
- **What it shows:**
 - Number of views over time.
 - Unique users viewing the report/dashboard.
 - Which content is most popular.

- Which users are most active.
 - **Purpose:** Helps **report owners** understand engagement (is my report being used? which pages/filters are most popular?).
 - **Limitations:**
 - Only available for reports/dashboards in a workspace.
 - Not a full audit trail (only high-level usage).
-

◆ 2. Audit Logs

- **What it is:** A more detailed, security-focused log of **user activities** across Power BI, available through **Microsoft 365 Compliance Center**.
 - **Scope:** Organization-wide (not just one report).
 - **Where to find it:**
 - Microsoft 365 Security & Compliance Center → **Audit Logs** → Search for Power BI activities.
 - Or via **PowerShell/Graph API**.
 - **What it shows:**
 - Report views, exports, sharing actions, dataset refreshes.
 - Who did what, when, and where (including IP/location).
 - Admin-level auditing for compliance & security.
 - **Purpose:** Helps **admins** monitor activity for governance, compliance, and security auditing.
 - **Requirements:**
 - Requires **Power BI Pro/Premium** and appropriate **Microsoft 365 auditing** enabled.
-

◆ Key Difference

- **Usage Metrics:** For **content creators** → “How often is my report/dashboard being used?”
 - **Audit Logs:** For **admins/compliance** → “Who accessed/exported/shared what, and when?”
-

✓ Example:

- A report author uses **Usage Metrics** to see if a sales dashboard is popular among regional managers.
- An IT admin uses **Audit Logs** to track whether someone exported sensitive data from that dashboard.

7. ◆ Steps to Manage Workspace Access

1. Go to the Workspace

- In Power BI Service → Left panel → **Workspaces** → Select the workspace you want.

2. Open Access Settings

- Click **Workspace settings** (gear icon) → **Permissions**.

3. Add People or Security Groups

- You can assign individual users, Microsoft 365 Groups, or security groups.
- Adding groups is recommended (easier to manage at scale).

4. Assign Roles (Permission Levels)

- Choose one of the predefined roles (see below).

◆ Workspace Roles & Permissions

Role	Can View Reports/Dashboards	Can Edit Content	Can Publish/Update Datasets	Manage Workspace Settings
Viewer	✓ Yes	✗ No	✗ No	✗ No
Contributor	✓ Yes	✓ Yes (edit/create)	✓ Yes	✗ No
Member	✓ Yes	✓ Yes	✓ Yes	✓ Limited (add users)
Admin	✓ Yes	✓ Yes	✓ Yes	✓ Full control

◆ Best Practices

- ✓ Use **security groups** instead of individuals to simplify management.

- Assign **Viewer** role for business users who just consume reports.
 - Assign **Contributor/Member** only to report developers or analysts.
 - Limit **Admin** to a few trusted people (workspace owners).
 - For publishing to a wide audience, use **Apps** (read-only packaged content) instead of giving direct workspace access.
-

Example:

- Sales managers = **Viewer** (they just consume dashboards).
- BI developers = **Contributor** (they can build and publish reports).
- IT admin = **Admin** (full workspace control).

8. ◆ 1. Access Control & Permissions

- **Workspaces roles** (Viewer, Contributor, Member, Admin) control who can view, edit, or manage content.
 - **Row-Level Security (RLS)** restricts access to specific rows of data based on user roles.
 - **Object-Level Security (OLS)** (Premium only) hides entire tables or columns.
 - **Conditional Access policies** (via Azure AD) restrict access based on device, location, or MFA.
-

◆ 2. Data Classification & Sensitivity Labels

- Apply **Microsoft Information Protection (MIP) sensitivity labels** to datasets, reports, and dashboards.
 - Labels travel with the data (Excel, PDF, etc.), so exports remain protected.
 - Example: “Confidential – Finance Only” → prevents unauthorized sharing.
-

◆ 3. Certified & Promoted Datasets

- **Promoted datasets** = recommended for broader use.
- **Certified datasets** = validated by data stewards.
- Helps enforce a **single source of truth** instead of multiple versions of the same data.

- ◆ **4. Monitoring & Auditing**

- **Usage Metrics** → Track report/dashboard usage at a high level.
 - **Audit Logs** (via Microsoft 365 Compliance Center) → Detailed activity trail (who viewed, shared, exported, or refreshed data).
 - **Admin portal** → Monitor dataset refresh failures, capacity usage, and permissions.
-

- ◆ **5. Dataflows & ETL Governance**

- Centralize transformation logic in **Dataflows** (instead of duplicated Power Query in multiple reports).
 - Store and manage **business rules once**, enforce consistency across reports.
-

- ◆ **6. Scheduled Refresh & Gateways**

- Use **scheduled refresh policies** to keep data accurate and up to date.
 - Manage on-premises data securely with **Gateways** and restrict who can publish or refresh datasets.
-

- ◆ **7. Capacity & Performance Management**

- **Premium Capacity** allows resource governance (memory, refresh concurrency, query limits).
 - **Deployment pipelines** enforce structured dev → test → prod content promotion.
-

- ◆ **8. Governance Policies & Training**

- Define clear **naming conventions** for workspaces, datasets, and reports.
 - Create **data stewardship roles** (owners responsible for certified data).
 - Educate users on **self-service BI best practices** while keeping guardrails in place.
-

 **Summary:**

Power BI Service enforces governance through a mix of **security (RLS,**

permissions), compliance (labels, audit logs), consistency (certified datasets, dataflows), and monitoring (usage, admin tools). This balance allows **self-service BI** while maintaining enterprise-level control.

9. ◆ Limitations of RLS with DirectQuery

1. Performance impact

- Every query must include the RLS filter conditions → can slow performance significantly, especially on large fact tables.
- Query folding depends on the data source; some sources handle filters poorly.

2. Limited transformations

- If your RLS logic relies on complex DAX or calculated tables, it may not fold into the SQL query → forcing Power BI to do extra work.

3. Cross-database scenarios

- If your model combines DirectQuery with Import (composite model), applying RLS across them can be tricky and may result in unsupported scenarios.

4. Data source limitations

- Not all DirectQuery sources support dynamic security well. For example:
 - SQL Server =  works.
 - Some sources (like SAP BW, Snowflake) = may have restrictions.

◆ Limitations of RLS with Live Connection

1. Defined at source, not Power BI

- With Live Connection (e.g., SSAS Tabular, Azure Analysis Services, Fabric semantic models), RLS is managed **only in the source model**, not inside Power BI.
- You cannot create or edit RLS roles in Power BI Service for these connections.

2. Dependency on IT/central modelers

- Business users cannot modify security rules → they must request changes from the central BI/IT team managing the semantic model.

3. Limited flexibility

- Since RLS lives in SSAS/Analysis Services, you can't combine it with Power BI's self-service features like calculated roles.
-

◆ General Limitations Across Both

- **Not applied to dataset owners/admins** → They always bypass RLS.
 - **Export/Analyze in Excel** → RLS applies, but sometimes users can infer restricted values if aggregation reveals patterns.
 - **Testing** → In Power BI Desktop, you must “View as Role” to test; in Service, you test via “Test as role” → harder in Live Connection.
 - **Licensing dependency** → Users consuming RLS-protected content must have **Power BI Pro or Premium per user (PPU)**.
-

✓ Summary:

- **DirectQuery:** RLS works, but can hurt performance and may not fold well to source queries.
- **Live Connection:** RLS must be managed at the source (SSAS/Analysis Services/Fabric), not in Power BI.
- **Overall:** RLS gives security but introduces trade-offs in performance, manageability, and flexibility.

10. ◆ 1. Refresh with Power Automate

Power Automate has a **built-in Power BI connector**.

Steps:

1. Go to Power Automate.
2. Create a new **flow** (automated, scheduled, or triggered).
3. Add the **Power BI → Refresh a dataset** action.
4. Select:
 - **Workspace** (where dataset lives)
 - **Dataset** (the one to refresh)
5. Add triggers if needed:
 - Schedule (e.g., every 2 hours).

- Event-based (e.g., when new file lands in SharePoint).
6. Save & test → The flow will call the refresh.
- Useful when: Business workflows (e.g., refresh dataset after ETL finishes or after user approval).
-

◆ 2. Refresh with Power BI REST API

For advanced or automated scenarios.

Endpoint:

POST

<https://api.powerbi.com/v1.0/myorg/groups/{groupId}/datasets/{datasetId}/refreshes>

Steps:

1. Register an app in **Azure AD** → get **Client ID** & permissions.
 - Permission needed: Dataset.ReadWrite.All.
 2. Authenticate via **OAuth2** to get an access token.
 3. Call the API with headers:
 4. Authorization: Bearer {access_token}
 5. Content-Type: application/json
 6. Example request body (optional):
 7. {
 8. "notifyOption": "MailOnFailure"
 9. }
 10. You can also check refresh status with:
 11. GET
<https://api.powerbi.com/v1.0/myorg/groups/{groupId}/datasets/{datasetId}/refreshes>
- Useful when: Integrating refresh into **CI/CD pipelines, custom apps**, or triggering refresh from external systems.
-

◆ 3. Key Differences

Feature	Power Automate	REST API
Ease of Use	✓ Low-code UI	✗ Requires coding
Flexibility	Medium	High (can embed in apps/scripts)
Triggers	Many (events, schedules)	Must be coded/scheduled separately
Monitoring	Limited	Full control (logs via API)

✓ Example Use Case

- **Power Automate:** Refresh sales dataset when a new Excel file is uploaded to OneDrive.
- **REST API:** Nightly ETL pipeline in Azure Data Factory calls Power BI API to refresh datasets once staging tables are loaded.